Language-Conditioned Pick-and-Place in Simulation: a Clearance-Aware Grasp Pipeline and a Study of Cross-Attention Pooling

Sattwik Singh
Department of Electrical Engineering, IIT Delhi
Code: `https://github.com/Sattwik-Singh/grasp-vla`

November 2026

Abstract

I re-implement, from scratch and in an open simulation, a language-conditioned pick-and-place system based on a Robotics Club project (GRASP) whose original code is no longer available. A Franka Panda in PyBullet parses an instruction, detects objects in an overhead RGB-D image, decides which one is meant, and grasps it with a clearance-aware planner. On 200 unseen scenes per condition, the pipeline succeeds in 99–100% of episodes across clean, cluttered, noisy-sensor, low-friction and combined conditions. Ablations show that splitting touching objects by colour removes almost all wrong-object errors (success 90.5%→99.0% in the hardest condition). A stress sweep shows the real weak point is camera calibration: success falls to 71% at 15 mm of error. Object spacing, even down to 1 mm, barely matters. In a controlled comparison of four ways to pool object tokens in a learned grounding policy, cross-attention pooling with an instruction-derived query cuts the correlation between representations of different instructions by 63% (0.951→0.348) and is the only variant that learns the task within the same 12-epoch budget (98.7% grounding accuracy against a 16.8% chance level; an instruction-modulated mean pool reaches 61% after 40 epochs). Inside the robot loop it completes 94% of cluttered episodes. I also report where the system fails and what this setup cannot show.

1 Introduction

This report rebuilds, from scratch, a language-conditioned grasping project (GRASP) that I worked on at the IIT Delhi Robotics Club. The original system combined open-vocabulary detection [3, 2], LLM-based task planning [4] and a grasp planner, and used cross-attention pooling in a vision-language-action (VLA) style policy [5, 6]. The original code and logs are no longer available to me, so its reported figures (an 85% reduction in cross-instruction correlation and 90% grasp success in simulation) cannot be checked. Instead, I re-implemented the two ideas behind those figures in an open, fully reproducible simulation and measured them again under controlled conditions. Two questions guide the work:

Q1 How reliable is a modular *instruction* → *detect* → *ground* → *grasp* pipeline, and where does it fail?

Q2 Does cross-attention pooling make a policy's internal representation instruction-specific, and does that translate into choosing the right object?

All numbers below come from scripts in the repository, run on fresh random scenes whose seeds were never used during development.

2 Simulation setup

Scene. A Franka Panda arm in PyBullet [1] stands at the edge of a table with a tray. Each scene places 5 or 8 objects at random positions and orientations. Each object is a cube, an upright cylinder or a ball (radius or half-size 2.0–2.8 cm, mass 50–200 g) in one of six colours, and no colour–shape pair appears twice in a scene. A fixed overhead camera gives a 320×320 RGB-D image. The simulator's segmentation masks and true poses are used only to *score* episodes; the robot never sees them.

Instructions. Instructions are drawn from templates with synonyms ("block", "can", "marble", ...), e.g. *"grab the blue tube and drop it in the basket"*. A quarter name two objects (*"move the green marble to the bin, then the red cylinder"*).

Conditions. Five conditions test different failure sources: *clean* (5 objects, at least 3 cm apart); *clutter* (8 objects, gaps down to 8 mm); *sensor noise* (2 mm depth noise plus an unknown camera offset with 6 mm

standard deviation, so the perception code works with a miscalibrated camera); *low friction* ($\mu \in [0.3, 0.6]$); and *combined* (all of the above, $\mu \in [0.4, 1.0]$).

Success. A step succeeds if the named object is lifted at least 10 cm, held for one simulated second, and ends up inside the tray. An episode succeeds only if every step succeeds. Failed steps are labelled *parse failure*, *no detection*, *wrong object*, *grasp miss*, *drop* or *place miss*.

Seeds. Development used scene seeds 0–39, a pilot on 1000–1029 (which exposed merged detections), and a first full run on 2000–2199 (whose grasp failures motivated the clearance-aware planner). All reported pick-and-place results use seeds 3000–3199, the stress sweep 5000–5099 and the learned-policy loop 4000–4099.

3 The modular pipeline

Planner. A rule-based parser turns the instruction into a list of *(colour, shape)* targets and decides whether to place them in the tray. The repository also contains a drop-in LLM planner that returns the same JSON plan, but it is not used in any reported result, so every number is reproducible offline.

Detection. The depth image is back-projected to world points. Points more than 6 mm above the table, outside the tray, are grouped into connected components. Touching objects merge into one component, so each component is split by 2-means clustering on pixel chromaticity when the two cluster centres are clearly apart (*colour splitting*, applied recursively). This step uses no colour names. Each detection yields a grasp point, the angle of its smallest enclosing rectangle and an RGB crop.

Grounding. The default grounder labels each detection's colour by nearest chromaticity and its shape from two geometric cues: how much of its bounding rectangle the footprint fills (squares ≈0.90, discs ≈0.79) and how domed its top is. The thresholds were tuned on the development scenes. An open-vocabulary CLIP [2] grounder that scores image crops against the phrase is implemented, but its weights could not be downloaded in the compute environment I used, so it is not evaluated here.

Grasp planning. The *simple* planner aligns the gripper with the object's rectangle and opens the fingers fully (4 cm each side). The *clearance-aware* planner opens the fingers only to the object's width plus 1.2 cm. It then tries candidate wrist angles (the two face-aligned angles for boxes, 15° steps for round objects) and keeps the angle whose finger footprints overlap the fewest points from neighbouring objects in the height map. The arm then moves in Cartesian straight lines using inverse kinematics.

4 Learned grounding and cross-attention pooling

To study the pooling question in isolation, I trained a small language-conditioned grounding policy. It sees one token per detected object (a CNN on the object's 24×24 RGB-D crop plus an MLP on its position, followed by LayerNorm). It encodes the instruction with a bidirectional GRU, and predicts the table position of the object to pick. The object tokens are pooled into one vector v, the fused representation is $h = \mathrm{MLP}([v; \ell])$, where ℓ is the language vector, and a linear head maps h to (x, y). Only the pooling step differs between the four variants:

- **Mean**: v is the average token, so it is the same for every instruction.
- **Attention (learned query)** [8]: attention with a fixed learned query. It has the same capacity as cross-attention, but its pooling still ignores the instruction.
- **FiLM + mean** [9]: every token is scaled and shifted by the instruction before mean pooling. This is a stronger instruction-aware baseline.
- **Cross-attention** [7]: the attention query is a projection of ℓ, so the instruction decides how much each object contributes.

Cross-instruction correlation. For a test scene with K objects, I encode the same scene under the K instructions that name each object, and average the pairwise Pearson correlations between the resulting fused vectors $h_1, \ldots, h_K$:

$$\rho = \frac{1}{K(K-1)} \sum_{i \neq j} \mathrm{corr}(h_i, h_j).$$

A value near 1 means the representation barely changes with the instruction, which suggests the policy is not really listening. I use this as my definition; the original project's exact definition is not available. Since a low correlation is only useful if the policy picks the right object, I also report *grounding accuracy*, the fraction of test instructions for which the object nearest the predicted point is the named one.

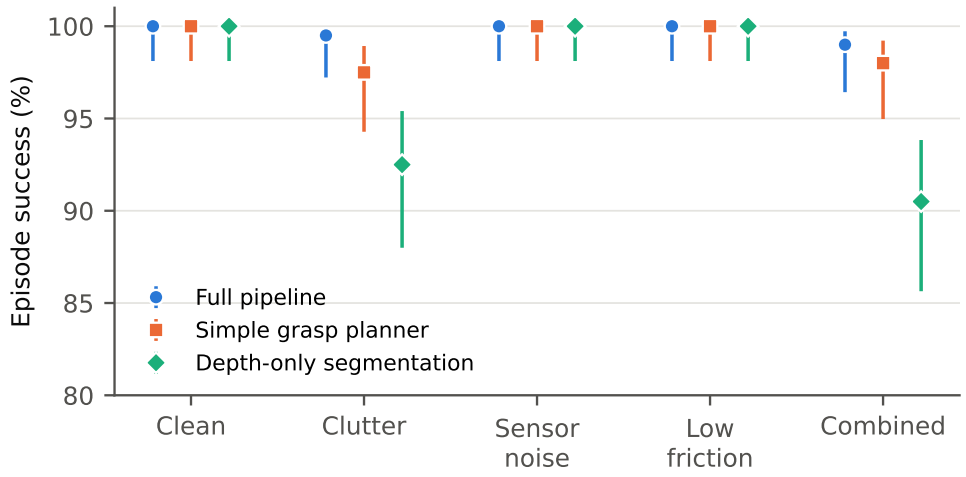


Figure 1: Episode success on 200 unseen scenes per condition (seeds 3000–3199), with 95% Wilson intervals. Blue: full pipeline. Orange: simple grasp planner. Green: depth-only segmentation (no colour splitting).

Condition	Full pipeline	Simple grasp planner	Depth-only segmentation
Clean	100%	100%	100%
Clutter	**99.5%** [97.2, 99.9]	97.5% [94.3, 98.9]	92.5% [88.0, 95.4]
Sensor noise	100%	100%	100%
Low friction	100%	100%	100%
Combined	**99.0%** [96.4, 99.7]	98.0% [95.0, 99.2]	90.5% [85.6, 93.8]
Failed steps in clutter (254 steps):			
wrong object / grasp miss	1 / 0	1 / 4	14 / 1

Table 1: Pick-and-place success (200 episodes per cell, 54 of them with two targets; 95% intervals in brackets). No episode failed through parsing, a missed detection, a drop or a missed placement.

Data and training. 2000 training scenes (12 070 instructions) and 300 test scenes (1 785 instructions), each with 4–8 objects and gaps of 8–30 mm. Each variant is trained for 12 epochs (AdamW, one-cycle learning rate peaking at 10^{-3}, MSE loss) with three seeds, on a 2-core CPU.

5 Results

5.1 Pick-and-place (Q1)

Figure 1 and Table 1 show that the full pipeline succeeds in 99–100% of episodes in every condition. Three findings stand out. (i) **Clutter is the only condition that causes failures**; sensor noise, camera miscalibration of a few millimetres and low friction cause none. (ii) **Colour splitting matters most.** Without it, touching objects merge into one detection and the robot picks the wrong one: wrong-object steps rise from 1 to 14 in clutter and from 1 to 18 in the combined condition, and success drops by 7–8.5 points. (iii) **The clearance-aware planner** reduces grasp misses in clutter (4 to 0) and in the combined condition (3 to 1), but the difference is inside the confidence intervals. Section 5.2 pushes spacing and calibration further to show where each version breaks.

5.2 Where does it break?

Figure 2 varies one difficulty knob at a time. **Spacing barely matters**: even with gaps of 1 mm between eight objects, success stays at 97%, and the few failures are wrong-object choices between touching objects of similar colour. **Camera calibration is the real weak point.** Success is 100% up to 5 mm of error, 96% at 10 mm, 71% at 15 mm and 29% at 30 mm, and almost every failure is a grasp miss. The system trusts the camera's geometry completely and has no feedback (such as a wrist camera or a re-grasp) to correct a misplaced gripper.

The two grasp planners behave almost identically in both sweeps. Pooling the two hardest conditions of Table 1, the clearance-aware planner had 1 grasp miss in 508 steps against 7 for the simple planner (Fisher's

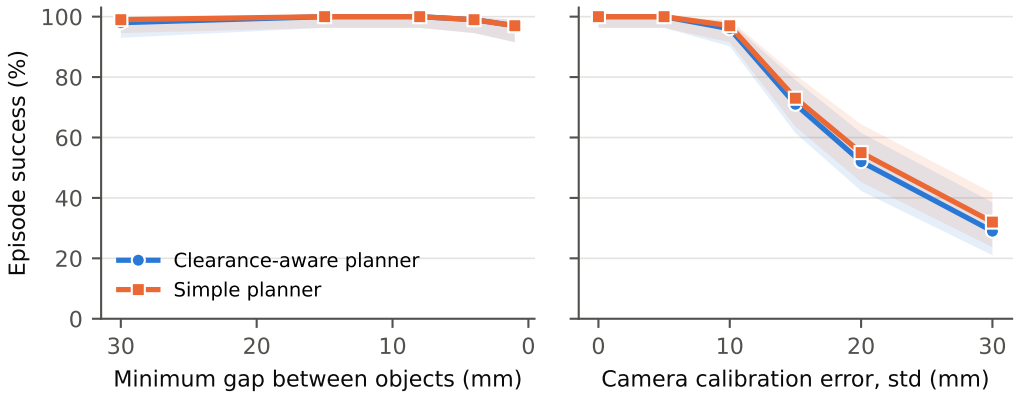


Figure 2: Stress sweep, 100 unseen scenes per point (seeds 5000–5099), with 95% intervals. Left: 8 objects packed ever more tightly. Right: 5 objects with an increasing unknown camera offset.

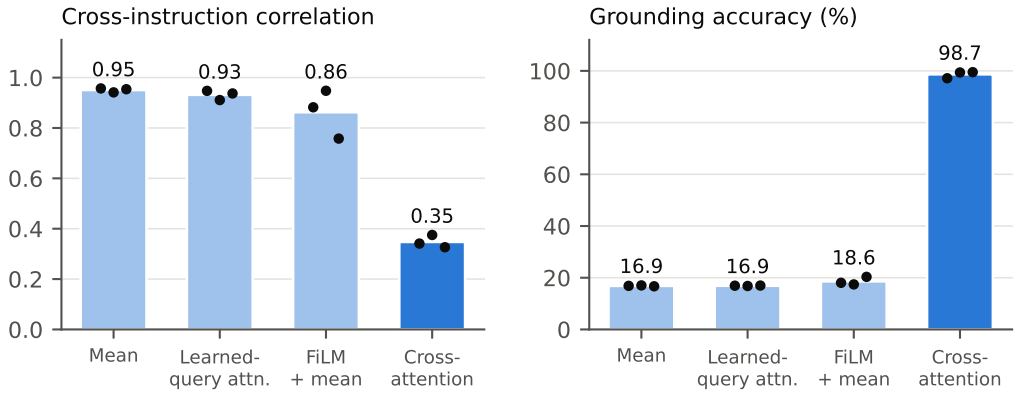


Figure 3: Held-out test scenes (300 scenes, 1 785 instructions). Bars show means over three training seeds; dots show individual seeds. Chance grounding accuracy is 16.8%.

exact test, $p = 0.07$). That is suggestive, not conclusive. Under large calibration error it is slightly *worse* (71/52/29% against 73/55/32% at 15/20/30 mm; not significant), which fits a real trade-off: opening the fingers less leaves less room for position error. An honest summary is that the planner fixes a specific clutter failure but does not measurably change overall reliability in this simulator.

5.3 Cross-attention pooling (Q2)

Cross-attention pooling lowers the cross-instruction correlation from 0.951 (mean pooling) to 0.348, a **63% reduction**. It is also the only variant that learns the task within the 12-epoch budget: 98.7% grounding accuracy and a 4.5 mm median error, against chance level (16.8%) for every other variant (Figure 3, Table 2). The learned-query attention baseline rules out "more parameters" or "attention in general" as the explanation: with a fixed query it behaves like mean pooling. What matters is that the *instruction* chooses what to attend to.

The FiLM baseline is also instruction-aware, yet it barely moves from chance within 12 epochs. Mean pooling forces it to represent "where the named object is" as an average over all objects, which is a much harder function to learn. To check whether this is just slow learning, I trained FiLM and mean pooling for 40 epochs (one seed each). FiLM does learn, slowly: 50.4% grounding accuracy after 20 epochs and 60.7% after 40, with a 49 mm median error, still far below cross-attention after 12 epochs. Mean pooling, by contrast, stays at chance (17.3%) after 40 epochs, even though its cross-instruction correlation falls to 0.59 and FiLM's to 0.43. So cross-attention is not the only pooling that can work, but it is far easier to learn: the instruction picks out one object directly, instead of the network having to cancel out every other object inside an average.

Two cautions. First, correlation and accuracy are not the same thing. After 40 epochs, mean pooling's correlation falls to 0.59 while its accuracy stays at chance, so a lower correlation alone is not evidence of better grounding. Second, attention weights are not a clean explanation of the decision. Figure 4 shows a success and a failure on the same scene. Pooled over three seeds, the error rate is highest when the target is a cylinder

Pooling	Cross-instr. correlation ↓	Grounding acc. ↑	Median error (mm) ↓
Mean	0.951 ± 0.008	16.9% ± 0.2	128.7
Attention (learned query)	0.932 ± 0.019	16.9% ± 0.1	129.0
FiLM + mean	0.863 ± 0.096	18.6% ± 1.5	125.4
Cross-attention	**0.348 ± 0.025**	**98.7% ± 1.3**	**4.5**

Table 2: Pooling comparison, mean ± standard deviation over three seeds, all trained for 12 epochs.

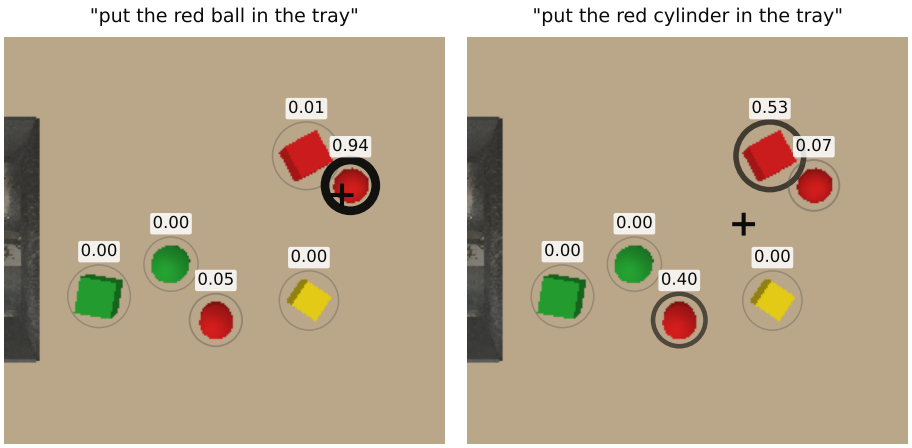


Figure 4: Cross-attention weights (numbers, ring thickness) and predicted grasp point (+) on one test scene. Left: the weight concentrates on the red ball and the prediction lands on it. Right: for the red cylinder (lower red object) the weight is split between the cylinder (0.40) and the red cube (0.53), and the prediction falls between them, a grounding failure.

(2.8%, against 0.9% for cubes and 0.3% for balls), because an upright cylinder seen from above looks like a ball or a cube.

5.4 The learned policy inside the robot loop

Replacing the hand-written grounder with each trained policy (seed-0 checkpoint) and running the full pick-and-place loop on 100 new cluttered scenes (seeds 4000–4099) gives the same ranking. Cross-attention succeeds in **94%** of episodes (95% interval 87.5–97.2) with 5 wrong-object steps. Mean, learned-query and FiLM pooling succeed in 8%, 8% and 6% of episodes, with 110 or more wrong-object steps each. The hand-written attribute grounder still does better (99%). In this closed world of six colours and three shapes, a few hand-written rules beat a learned model. The learned policy's advantage would show only with object descriptions the rules cannot express.

6 Relation to the original project

My CV lists two GRASP results, and this study re-measures both under definitions I state explicitly. *90% grasp success in simulation:* the re-implemented pipeline reaches 99–100% across five conditions on primitive objects. That number is higher, but it comes from a simpler simulator and is not comparable to the original. *85% reduction in cross-instruction correlation through cross-attention pooling:* with the definition in Section 4, I measure a 63% reduction (0.951 to 0.348), together with a jump in grounding accuracy from chance to 98.7%. The direction of both original results holds, but the specific numbers belong to the original setup and should not be quoted as reproduced.

7 Limitations

- **Easy world.** The objects are rigid primitives, grasps are top-down only, and the camera sees everything from above. High success here says little about real objects, occlusion or a real robot.
- **Closed vocabulary.** The default grounder knows six colours and three shapes, and its thresholds were

tuned by hand on development scenes. The CLIP grounder is implemented but could not be evaluated, and the LLM planner is optional and unevaluated. The "open-vocabulary" and "LLM planning" parts of the original project are therefore not tested here.

- **Small learned model.** The grounding policy uses the pipeline's own detections as tokens and is trained for 12 epochs. Baselines were compared at equal training budget. Given more epochs, FiLM pooling partly catches up (see Section 5.3).
- **Metric.** Cross-instruction correlation is my own definition. As the FiLM result shows, it should always be reported together with a task metric.

8 Conclusion

A simple modular pipeline with two targeted fixes, colour splitting of merged detections and clearance-aware finger placement, picks and places named objects in 99–100% of unseen simulated scenes. Its main weakness is its complete trust in camera calibration: beyond about 10 mm of error, grasps start to miss. In a controlled comparison, letting the instruction form the attention query made a small policy learn to ground instructions quickly and accurately. It lowers the correlation between representations of different instructions by 63% and lifts grounding accuracy from chance to 98.7%. The natural next steps are to evaluate the open-vocabulary CLIP grounder and LLM planner, add realistic object meshes and side-on grasps, and test the learned policy on descriptions the hand-written rules cannot express.

Acknowledgements and AI use. The original GRASP project was a team effort at the IIT Delhi Robotics Club. This report is my independent re-implementation. The code and this report were developed with AI assistance (Claude). I directed the work and take responsibility for the results.

References

[1] E. Coumans, Y. Bai. PyBullet, a Python module for physics simulation for games, robotics and machine learning. `http://pybullet.org`, 2016–2021.

[2] A. Radford et al. Learning transferable visual models from natural language supervision. *ICML*, 2021.

[3] M. Minderer et al. Simple open-vocabulary object detection with vision transformers. *ECCV*, 2022.

[4] M. Ahn et al. Do as I can, not as I say: Grounding language in robotic affordances. *CoRL*, 2022.

[5] M. Shridhar, L. Manuelli, D. Fox. CLIPort: What and where pathways for robotic manipulation. *CoRL*, 2021.

[6] A. Brohan et al. RT-2: Vision-language-action models transfer web knowledge to robotic control. *CoRL*, 2023.

[7] A. Vaswani et al. Attention is all you need. *NeurIPS*, 2017.

[8] J. Lee et al. Set Transformer: A framework for attention-based permutation-invariant neural networks. *ICML*, 2019.

[9] E. Perez et al. FiLM: Visual reasoning with a general conditioning layer. *AAAI*, 2018.